Cache Optimisations and CSR Sparse Storage for Dense Matrix Multiplication: A Controlled Benchmarking Study

Alejandro Delgado Valera
BIG DATA – Individual Assignment 2
Grado en Ciencia e Ingeniería de Datos
Universidad de Las Palmas de Gran Canaria
https://github.com/aledelgadoo/individual-assignments-bigdata

October 20, 2026

Abstract

Questions: (A) how do loop reordering and cache blocking affect the runtime and memory of dense $O(n^3)$ multiplication relative to a validated triple-loop baseline, and (B) under which tested sparsity levels does a CSR sparse representation beat dense methods on matched logical inputs, once conversion costs are measured separately? All four methods are single-threaded C kernels (`float64`, flags `-O2 -std=c11`) with bitwise-identical results, measured on the multiplication kernel with warm-ups, five batched repetitions and a declared budget of 60 s and 4 GiB. Loop reordering to `ikj` yields 2.5–2.8× over the `ijk` baseline at $n \in \{500, 1000\}$ and 7.0× at $n = 2000$; cache blocking with tile size 64 is slower than `ikj` up to $n = 1000$ but 1.6× faster at $n = 2000$, where the 32 MB B matrix no longer fits in the 12 MB last-level cache. Dense runtimes are flat across sparsity levels; CSR time falls monotonically from 263 ms (1 % zeros) to 4.0 ms (99 % zeros) at $n = 1000$, beating every dense method on the whole tested grid, while dense-to-CSR conversion (1.4–10.7 ms) amortises within one multiplication. Peak resident memory grows from 26.7 MB (dense) to 38.1 MB (CSR) at $n = 1000$. Conclusions are scoped to these kernels, this machine and uniformly random sparsity.

1 Introduction and objective

Assignment 1 established the baseline for this study: a straightforward i–j–k dense multiplication whose constant factors are dominated by memory behaviour rather than by the $2n^3$ arithmetic once the matrices leave cache [1, 2]. This assignment asks which cheap, purely local changes improve that baseline, and whether a sparse storage format pays off when most entries of A are zero.

Research questions.

A: How do loop reordering (`ikj`) and cache blocking (`block`) affect dense multiplication time and memory relative to the validated `ijk` baseline, under one controlled protocol?

B: Under which tested sparsity levels does a CSR sparse implementation outperform the dense methods on matched logical inputs, with conversion costs measured separately?

Hypotheses (stated before measurement).

H1: `ikj` is faster than `ijk` because its inner loop streams B by rows instead of by strided columns; the advantage grows with n as strided access stresses the TLB and caches.

H2: Blocking pays off only when the working set exceeds the last-level cache; at small and medium n its bookkeeping makes it no faster than `ikj`.

H3: The dense methods do not get faster with sparsity: skipping zeros is not what the dense loops do.

H4: CSR time tracks the number of nonzeros, so it decreases with sparsity and crosses below the best dense method at some tested level.

H5: All four methods produce bitwise-identical results on matched inputs, because they share one summation order.

The controlled factor is the method; the outcome measures are kernel time per multiplication, peak resident memory, and (for CSR) conversion time.

2 Background

Loop order and memory access. In row-major storage the baseline `ijk` loop reads B down columns with stride $8n$ bytes, which destroys spatial locality and stresses the TLB at large n. The reordered `ikj` loop traverses B by rows (contiguous 64-byte lines) and accumulates into a contiguous row of C, trading a read-modify-write of C for streaming accesses everywhere [3, 2]. Both perform the same $2n^3$ FLOPs in the same k order, so their floating-point results are bitwise equal [7].

Cache blocking. Blocking partitions the k loop into tiles of width B so that the $B \times B$ tiles of A, B and C stay resident while reused. The tile working set is roughly $3 \times 8B^2$ bytes: 24 KB at $B = 32$ (fits the 48 KB L1d) and 98 KB at $B = 64$ (fits the 1.25 MB L2). The optimal tile is a compromise between those levels; we select it empirically instead of assuming one value [3].

Compressed sparse row (CSR). CSR stores the nonzeros of A in three arrays: `rowptr` (n+1 indices), `colind` (nnz column indices) and `vals` (nnz values). Multiplying $C = AB$ then loops over the nonzeros of each row of A and streams the corresponding dense rows of B and C [4]. The arithmetic work becomes $2 \cdot nnz \cdot n$ instead of $2n^3$, at the price of storage that can exceed $8n^2$ when A is dense and of a one-off conversion from the dense input.

Benchmarking practice. We keep the median/IQR summary, the raw measurements, the single timing boundary and the controlled environment of Assignment 1 [5, 6]. One additional laptop-specific control appears in Section 3: the CPU frequency policy (Section 5).

3 Methodology

3.1 Implementations

Four single-threaded C11 kernels in one binary (`src/matmul2.exe`), compiled with the fixed flags `-O2 -std=c11 -Wall -Wextra` (GCC 16.1.0, MinGW-w64); no numerical library runs inside the kernel and all four are sequential:

- `ijk` – the Assignment 1 baseline kernel, reused unchanged from Git tag `assignment-1` (`src/c/matmul.c`); validated in that study against NumPy and the other languages.
- `ikj` – same arithmetic, reordered loops so the inner `j` loop streams B by rows; `restrict` pointers.
- `block` – tiled i/j loops over k-tiles of width $B \in \{16, 32, 64, 128\}$, selected per n by a tuning sweep (Section 4); 64 won at every tuned size.
- `csr` – dense-to-CSR conversion (two-pass with a cursor array), then CSR×dense: for each row of A, each nonzero `(k, val)` streams row k of B into row i of C.

Every kernel zeroes C inside its timed region (the operation includes writing the output), uses `float64` arithmetic, and visits k in ascending order, so all four emit bitwise-identical results on matched inputs (verified in Section 3.3). Input loading, allocation, conversion (CSR) and checksums are outside the timed region. The process is pinned to logical CPU0 (a P-core of the hybrid i7-1355U) with `SetThreadAffinityMask`.

3.2 Inputs

Dense Part A reuses the Assignment 1 input pairs byte-for-byte. For each declared size $n = 10$, 50, 100, 200, 500, 1000 and 2000, A and B were generated with NumPy PCG64 seed scheme `default_rng([42, n, index])`, uniform in $[-1, 1)$, stored as little-endian `<f8` row-major files, with SHA-256 checksums in `data/inputs/meta.json`.

Sparse Part B keeps $n = 1000$ and B dense, and varies only A: nine sparsity levels $p \in \{1, 5, 10, 25, 50, 75, 90, 95, 99\}$ % of zeros, produced by masking the dense A_{1000} with an independent keep-mask (`default_rng([42, n, p, 2])`, keep if uniform $\geq p/100$), so the nonzero distribution is uniform random and the same file feeds all four methods. Recorded nonzero counts match the targets (e.g. $nnz = 499\,827$ at 50 %, 9 844 at 99 %).

3.3 Correctness

Before any benchmarking, the self-test executable (`matmul2.exe selftest`) ran 304 checks, all passing: sizes 1–64 crossed with four matrix patterns (dense random, 50 % zeros, zero matrix, identity) and four block sizes $\{4, 7, 16, 64\}$; each method was compared against a `long double` reference ($rtol = 10^{-10}$, $atol = 10^{-12}$, tight because the summation orders are identical), bitwise between methods, plus manual nnz counts and hand-computed $2 \times 2/3 \times 3$ answers. Independently, the campaign CSVs show one unique checksum per matched input across the four methods for all 16 (n, p) groups.

3.4 Environment

CPU	13th Gen Intel Core i7-1355U, 10 physical / 12 logical processors
Cache	L1d 48 KB per P-core, L2 6656 KB total, L3 12288 KB, line 64 B
RAM / OS	16.8 GB, Windows 10 (build 26200)
Compiler	GCC 16.1.0 (MinGW-w64), `-O2 -std=c11 -Wall -Wextra`
Power scheme	Windows *High performance* (turbo clocks held)
Affinity	children pinned to logical CPU0 via `SetThreadAffinityMask`

The first campaign ran under the default *Balanced* plan and was discarded: single-core frequency oscillated between turbo and throttled states, producing 5× swings inside one configuration. All reported numbers come from a re-run under *High performance*; the effect itself is discussed in Section 5.

3.5 Measurement protocol

Identical to Assignment 1, with one binary: monotonic `QueryPerformanceCounter` timing of the kernel only; a feasibility probe (one excluded warm-up + one timed kernel, retried once, gate = timed $\leq$ 60 s); then, for passing configurations, one excluded warm-up followed by five measured repetitions, batched so each timing spans at least 0.2 s (batch sizes calibrated at run time; very slow kernels use batch = 1). Summaries are the median and IQR of the five per-multiplication times. Memory is the peak resident (working-set) size of the child process, sampled every 50 ms by the orchestrator (psutil on Windows) and by the process itself (`GetProcessMemoryInfo`); budget 4 GiB. CSR conversion is timed

Table 1: Part A: median kernel time per multiplication (ms) over 5 repetitions; IQR as a percentage of the median ranged 0.2–17.8 % (median 4.3 %).

n	`ijk`	`ikj`	`block` ($B = 64$)
10	0.0004	0.0005	0.0005
50	0.033	0.023	0.039
100	0.300	0.213	0.340
200	3.22	1.84	2.85
500	75.07	29.55	41.54
1000	806.9	286.1	325.2
2000	27 478	3 941	2 525

Table 2: Tuning sweep: median kernel time (ms) for four candidate block sizes.

n	$B = 16$	$B = 32$	$B = 64$	$B = 128$
500	80.8	85.1	**46.7**	57.7
1000	960	596	**313**	430
2000	11 377	7930	**2627**	3691

in its own CSV rows (`conversions.csv`), never inside the multiplication timing. System CPU utilisation is recorded per configuration as a noise covariate. Raw data live in `data/results/{raw_timings,configs,conversions,blocksize_tuning}.csv`.

4 Results

All 57 configurations (21 dense + 36 sparse) completed within budget; there are no timeouts to report. Across configurations the IQR is a median of 4.3 % of the median (maximum 17.8 %, on the sub-millisecond $n = 200$–1000 configurations).

4.1 Part A: dense optimisation

Table 1 and Figure 1 show the scaling. `ikj` is faster than the baseline from $n = 50$ upward: 2.54× at $n = 500$, 2.82× at $n = 1000$ and 6.97× at $n = 2000$. Blocking follows a different curve: it is slower than `ikj` at every size up to $n = 1000$ (e.g. 1.41× slower at $n = 500$) but overtakes it at $n = 2000$, where it is 1.56× faster than `ikj` and 10.9× faster than the baseline (Figure 2). The local growth exponent between $n = 500$ and $n = 2000$ is 4.26 for `ijk`, 3.53 for `ikj` and 2.96 for `block`: the baseline is strongly super-cubic, and blocking is the only method close to the algorithmic n^3 at the top end.

The block size was chosen by a tuning sweep (Table 2); $B = 64$ won at every tuned size and was used in the campaign.

4.2 Part B: sparse multiplication at n = 1000

Table 3 and Figure 4 confirm H3 and H4: the three dense methods vary by less than 12 % across the whole sparsity range (their loops still touch every entry), while CSR falls monotonically from 263 ms to 4.0 ms. On this grid CSR is the fastest method at every sparsity level, including 1 %, where it is already 1.1× faster than `ikj`; the advantage grows to 25–90× at 90–99 % sparsity. The 2.4 ms–10.7 ms dense-to-CSR conversion (Figure 5) amortises after 0.07 multiplications at 50 % sparsity and after 0.01 at 90 %.

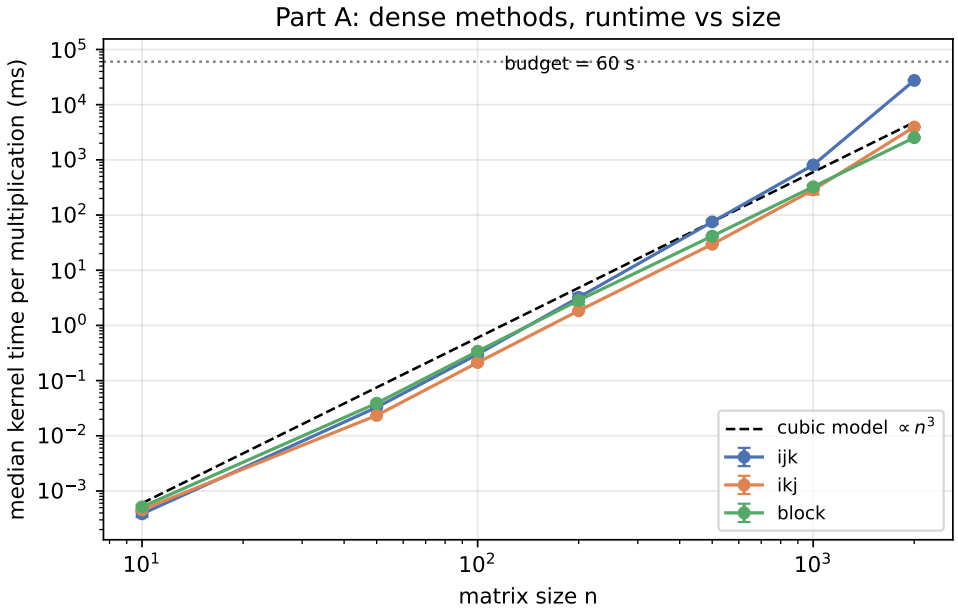


Figure 1: Part A: median kernel time vs. n (log-log) with IQR error bars and a cubic reference anchored at the baseline ($n = 500$). All configurations met the 60 s budget.

Table 3: Part B at $n = 1000$: median kernel time (ms) per method and measured nonzero count of A. The dense methods are flat in p; CSR tracks nnz.

sparsity	nnz	`ijk`	`ikj`	`block`	`csr`
1 %	989 972	802	296	314	**263**
5 %	949 910	817	317	324	**258**
10 %	900 120	813	271	322	**250**
25 %	749 936	865	281	323	**215**
50 %	499 827	773	298	335	**145**
75 %	249 722	807	284	317	**70.7**
90 %	100 375	782	313	312	**33.4**
95 %	50 189	795	326	357	**14.9**
99 %	9 844	810	357	348	**4.0**

4.3 Memory

Peak resident memory at $n = 1000$ is 26.7 MB for the dense methods (theoretical raw storage $3 \times 8n^2 = 24$ MB plus runtime overhead) and up to 38.1 MB with CSR: the sparse format adds the index and value arrays ($\approx 3 \times 4\,nnz$ bytes, plus the row pointer), i.e. about 11 MB at 99 % density and proportionally less as sparsity grows. No configuration came near the 4 GiB budget; the largest configuration ($n = 2000$, `ijk`, 27.5 s) stayed inside the 60 s budget, so the largest feasible tested size under the budget is $n = 2000$ for all three dense methods.

5 Discussion

Part A. H1 holds: reordering to `ikj` buys 2.5–2.8× at medium sizes and 7× at $n = 2000$, and the baseline's super-cubic local exponent (4.26) is consistent with strided column reads of B leaving cache and TLB reach (B is 32 MB at $n = 2000$ against a 12 MB L3). H2 holds in the form observed: blocking is pure overhead relative to `ikj` while the tiles fit in cache anyway (up to $n = 1000$), and only pays at $n = 2000$, where the k-tiling keeps 98 KB working sets resident

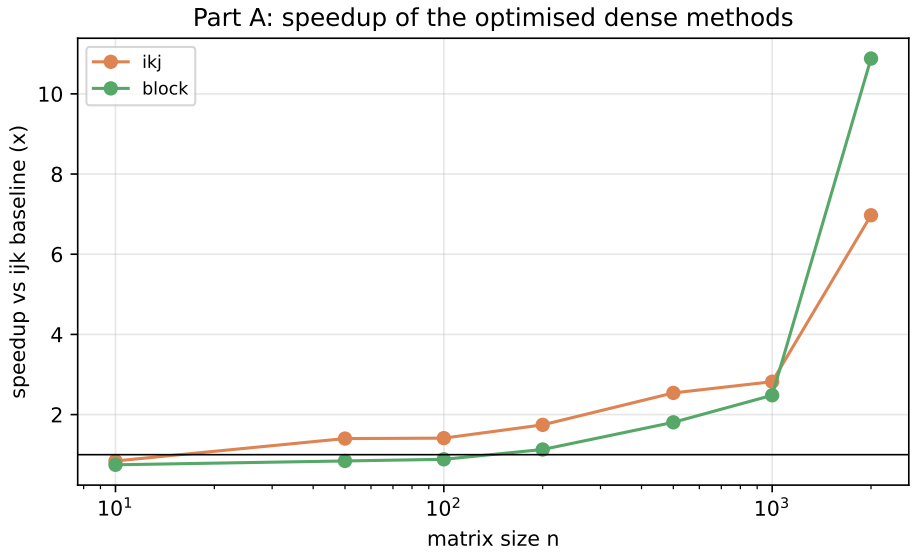


Figure 2: Part A: speedup of the optimised methods over `ijk` (median times). Blocking only pays off past $n = 1000$.

and restores near-cubic growth (2.96). The tuning sweep supports the cache argument: $B = 64$, whose three tiles fit the 1.25 MB L2, beat the L1d-fitting $B = 32$ at every tuned size, and both large candidates beat the too-small $B = 16$.

Part B. H3 holds: on this machine the dense loops gain nothing from zeros (a zero entry still costs a multiply and an add). H4 holds, with a qualification: on the whole tested grid CSR wins already at 1 % sparsity, not only at high sparsity. The reason visible in the kernel is locality, not just arithmetic: CSR reads A sequentially and streams B and C by rows, while even the best dense method (`ikj`) still pays an $A[i][k]$ load and a strided k step per output element. The conversion that CSR requires is cheap (1.4–10.7 ms) and amortises within a single multiplication, so it never erases the advantage at $n = 1000$. H5 holds by construction and by measurement: all checksums are bitwise identical.

Methodological finding. The discarded first campaign is itself a result: on a hybrid laptop, the *Balanced* power plan lets single-core frequency swing by factors of 2–5 within one configuration, which would have corrupted every cross-method comparison. Pinning and the *High performance* plan (or an equivalent fixed-frequency control) should be considered part of the protocol on this class of machine.

Limitations.

1. One machine, one compiler, sequential kernels only; the hybrid P/E core design means other laptops may rank the methods differently.

2. Part B tests a single size ($n = 1000$); the dense-vs-CSR crossover is size-dependent and should not be extrapolated to $n = 2000$ (CSR's memory traffic grows as $nnz \cdot n$, while blocked dense reuses cache).

3. Sparsity is uniform random; clustered or structured sparsity can change CSR's row-length distribution and its performance. Percentage of zeros alone does not determine performance.

4. All methods were auto-vectorised by the same compiler flags; no explicit vectorisation study was performed.

5. Memory differences are working-set metrics of one process, not allocator-level breakdowns.

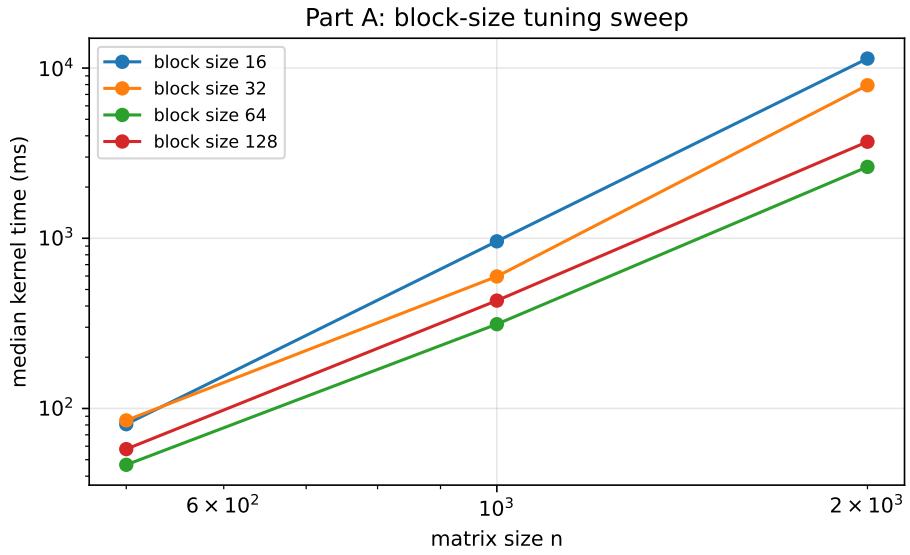


Figure 3: Tuning sweep of the block size. The 98 KB working set of $B = 64$ (3 tiles of 64×64 doubles) matches the per-core L2 better than the L1d-fitting $B = 32$.

Following the assignment guidance, the claims are limited to these implementations, inputs and hardware [5, 6].

6 Conclusions

Within the tested scope: (1) loop reordering to `ikj` is the reliable winning optimisation for dense multiplication on this CPU, worth 2.5–7× over the validated `ijk` baseline; (2) cache blocking is not free — it loses to `ikj` up to $n = 1000$ and wins only at $n = 2000$, where the inputs outgrow the last-level cache, with tile size 64 empirically best; (3) dense runtimes are insensitive to sparsity, while CSR time tracks the nonzero count and beats every dense method across all tested sparsity levels at $n = 1000$, from 1 % to 99 % zeros; (4) CSR's conversion cost is negligible (< 0.1 multiplications to amortise) and its memory overhead is about 11 MB at $n = 1000$; and (5) power management, not algorithmics, was the dominant noise source in the first campaign and had to be controlled explicitly.

Unresolved: the size dependence of the sparse crossover, structured sparsity, and whether blocking and CSR compose (blocked CSR) or interact with the parallel and distributed variants of Assignments 3 and 4.

7 Reproducibility materials

All source code, shared inputs, raw measurements (CSV) and the executed analysis notebook live in the public repository `https://github.com/aledelgadoo/individual-assignments-bigdata`; the Git tag `assignment-2` marks the exact state submitted for this assignment (the Assignment 1 baseline is preserved under tag `assignment-1`). The `README.md` documents dependency versions and the exact build, input-generation, benchmark, environment and analysis commands that regenerate this report end to end.

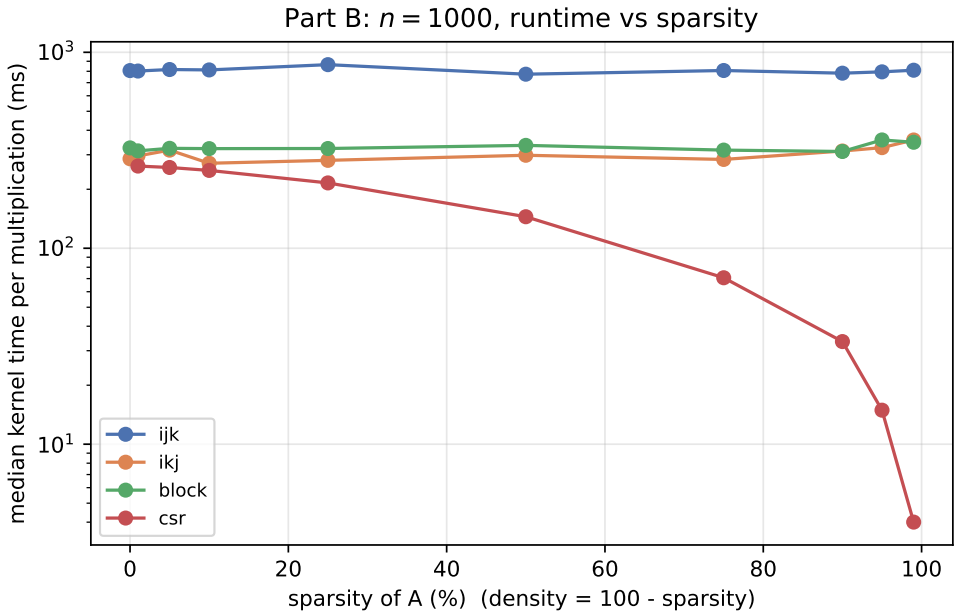


Figure 4: Part B: median kernel time vs. sparsity at $n = 1000$ (log y-axis). Dense methods are flat; CSR follows the nonzero count.

Generative-AI usage statement

Following the course policy, the author discloses that the AI assistant **opencode** (model `opencode/mimo-v2.6-flash-free`) was used, under the author's direction, for planning the optimised kernels and the measurement protocol, drafting and debugging parts of the C code and Python scripts, generating the analysis notebook and figures, and revising the English of this report. The author reviewed all generated code, executed and inspected the correctness self-test, the tuning sweep, the benchmark campaign and every figure, verified the numbers quoted here against the raw CSV files, and is responsible for the measurements, the conclusions and any errors.

References

[1] G. Strang, *Introduction to Linear Algebra*, 5th ed. Wellesley–Cambridge Press, 2016.

[2] T. H. Cormen, C. E. Leiserson, R. L. Rivest, and C. Stein, *Introduction to Algorithms*, 4th ed. MIT Press, 2022.

[3] G. H. Golub and C. F. Van Loan, *Matrix Computations*, 4th ed. JHU Press, 2013.

[4] Y. Saad, *Iterative Methods for Sparse Linear Systems*, 2nd ed. SIAM, 2003.

[5] W. P. Fleming and J. M. Wallace, "How not to lie with statistics: the correct way to summarize benchmark results," *Communications of the ACM*, vol. 29, no. 3, pp. 218–221, 1986.

[6] T. Mytkowicz, A. Diwan, M. W. Hauswirth, and P. F. Sweeney, *Measuring Computer Performance: A Practitioner's Guide*. Cambridge University Press, 2009.

[7] ISO/IEC 9899:2011, "Information technology — Programming languages — C," International Organization for Standardization, 2011.

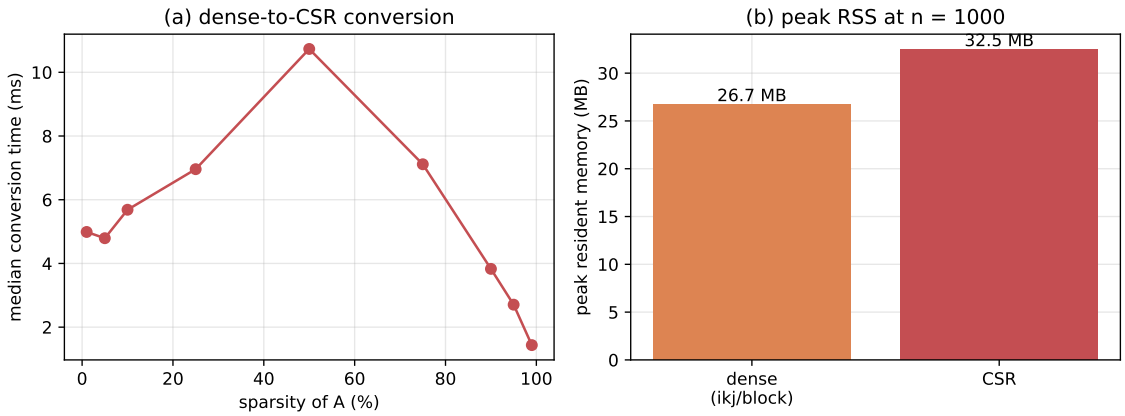


Figure 5: Part B: (a) median dense-to-CSR conversion time, measured separately; (b) mean peak resident memory of the child at $n = 1000$ (dense methods vs. CSR).

[8] G. Rodola, *psutil: cross-platform process and system utilities*, 2026. [Online]. Available: `https://github.com/giampaolo/psutil`